

ENV 790.30 - Time Series Analysis for Energy Data | Spring 2024

Assignment 5 - Due date 02/13/24

Vincient Whatley

Directions

You should open the .rmd file corresponding to this assignment on RStudio. The file is available on our class repository on Github. And to do so you will need to fork our repository and link it to your RStudio.

Once you have the file open on your local machine the first thing you will do is rename the file such that it includes your first and last name (e.g., “LuanaLima_TSA_A05_Sp23.Rmd”). Then change “Student Name” on line 4 with your name.

Then you will start working through the assignment by **creating code and output** that answer each question. Be sure to use this assignment document. Your report should contain the answer to each question and any plots/tables you obtained (when applicable).

When you have completed the assignment, **Knit** the text and code into a single PDF file. Submit this pdf using Sakai.

R packages needed for this assignment: “readxl”, “ggplot2”, “forecast”, “tseries”, and “Kendall”. Install these packages, if you haven’t done yet. Do not forget to load them before running your script, since they are NOT default packages.\

```
#Load/install required package here  
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':  
##   method      from  
##   as.zoo.data.frame zoo
```

```
library(tseries)  
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(tidyr)
```

```
## Warning: package 'tidyr' was built under R version 4.3.2
```

```
library(ggplot2)
library(Kendall)
library(lubridate)
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union
```

```
library(tidyverse) #load this package so you clean the data frame using pipes
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v forcats 1.0.0      v stringr 1.5.1
## v purrr  1.0.2      v tibble  3.2.1
## v readr   2.1.5
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library("xlsx")
```

Decomposing Time Series

Consider the same data you used for A04 from the spreadsheet “Table_10.1_Renewable_Energy_Production_and_Consumption”. The data comes from the US Energy Information and Administration and corresponds to the December 2023 Monthly Energy Review.

```
#Importing data set - using xlsx package
energy_data <- read.xlsx(file="./Data/Table_10.1_Renewable_Energy_Production_and_Consumption_by_Source.xlsx")
```

```
#Now let's extract the column names from row 11 only
read_col_names <- read.xlsx("./Data/Table_10.1_Renewable_Energy_Production_and_Consumption_by_Source.xlsx", sheet = 1, startRow = 11, endRow = 11, colNames = TRUE)
```

```
colnames(energy_data) <- read_col_names
head(energy_data)
```

```
##           Month Wood Energy Production Biofuels Production
## 1 1973-01-01                129.630             Not Available
## 2 1973-02-01                117.194             Not Available
```

## 3	1973-03-01	129.763	Not Available
## 4	1973-04-01	125.462	Not Available
## 5	1973-05-01	129.624	Not Available
## 6	1973-06-01	125.435	Not Available
##	Total Biomass Energy Production	Total Renewable Energy Production	
## 1		129.787	219.839
## 2		117.338	197.330
## 3		129.938	218.686
## 4		125.636	209.330
## 5		129.834	215.982
## 6		125.611	208.249
##	Hydroelectric Power Consumption	Geothermal Energy Consumption	
## 1		89.562	0.490
## 2		79.544	0.448
## 3		88.284	0.464
## 4		83.152	0.542
## 5		85.643	0.505
## 6		82.060	0.579
##	Solar Energy Consumption	Wind Energy Consumption	Wood Energy Consumption
## 1	Not Available	Not Available	129.630
## 2	Not Available	Not Available	117.194
## 3	Not Available	Not Available	129.763
## 4	Not Available	Not Available	125.462
## 5	Not Available	Not Available	129.624
## 6	Not Available	Not Available	125.435
##	Waste Energy Consumption	Biofuels Consumption	
## 1	0.157	Not Available	
## 2	0.144	Not Available	
## 3	0.176	Not Available	
## 4	0.174	Not Available	
## 5	0.210	Not Available	
## 6	0.176	Not Available	
##	Total Biomass Energy Consumption	Total Renewable Energy Consumption	
## 1		129.787	219.839
## 2		117.338	197.330
## 3		129.938	218.686
## 4		125.636	209.330
## 5		129.834	215.982
## 6		125.611	208.249

```
nobs=nrow(energy_data)
nvar=ncol(energy_data)
```

Q1

For this assignment you will work only with the following columns: Solar Energy Consumption and Wind Energy Consumption. Create a data frame structure with these two time series only and the Date column. Drop the rows with *Not Available* and convert the columns to numeric. You can use filtering to eliminate the initial rows or convert to numeric and then use the `drop_na()` function. If you are familiar with pipes for data wrangling, try using it!

```

#Filter for Solar and Wind Collumms
#Removing unnecessary columns, renaming columns
energy_data <- energy_data[,c(1, 8, 9)]
colnames(energy_data)=c("Date", "Solar Energy Consumption", "Wind Energy Consumption")

#Convert columns to numeric and replace "Not Available" with NA
energy_data <- energy_data %>%
  mutate(`Solar Energy Consumption` = if_else(`Solar Energy Consumption` == "Not Available", NA_real_,
    `Wind Energy Consumption` = if_else(`Wind Energy Consumption`=="Not Available", NA_real_,
      drop_na())

```

```

## Warning: There were 2 warnings in 'mutate()'.
## The first warning was:
## i In argument: 'Solar Energy Consumption = if_else(...)'.
## Caused by warning in 'if_else()':
## ! NAs introduced by coercion
## i Run 'dplyr::last_dplyr_warnings()' to see the 1 remaining warning.

```

```

#Converting to time series objects
energy_data$`Solar Energy Consumption` <- ts(energy_data$`Solar Energy Consumption`, frequency = 1)
energy_data$`Wind Energy Consumption` <- ts(energy_data$`Wind Energy Consumption`, frequency = 1)

```

Q2

Plot the Solar and Wind energy consumption over time using ggplot. Plot each series on a separate graph. No need to add legend. Add informative names to the y axis using `ylab()`. Explore the function `scale_x_date()` on ggplot and see if you can change the x axis to improve your plot. Hint: use `scale_x_date(date_breaks = "5 years", date_labels = "%Y")`)

```

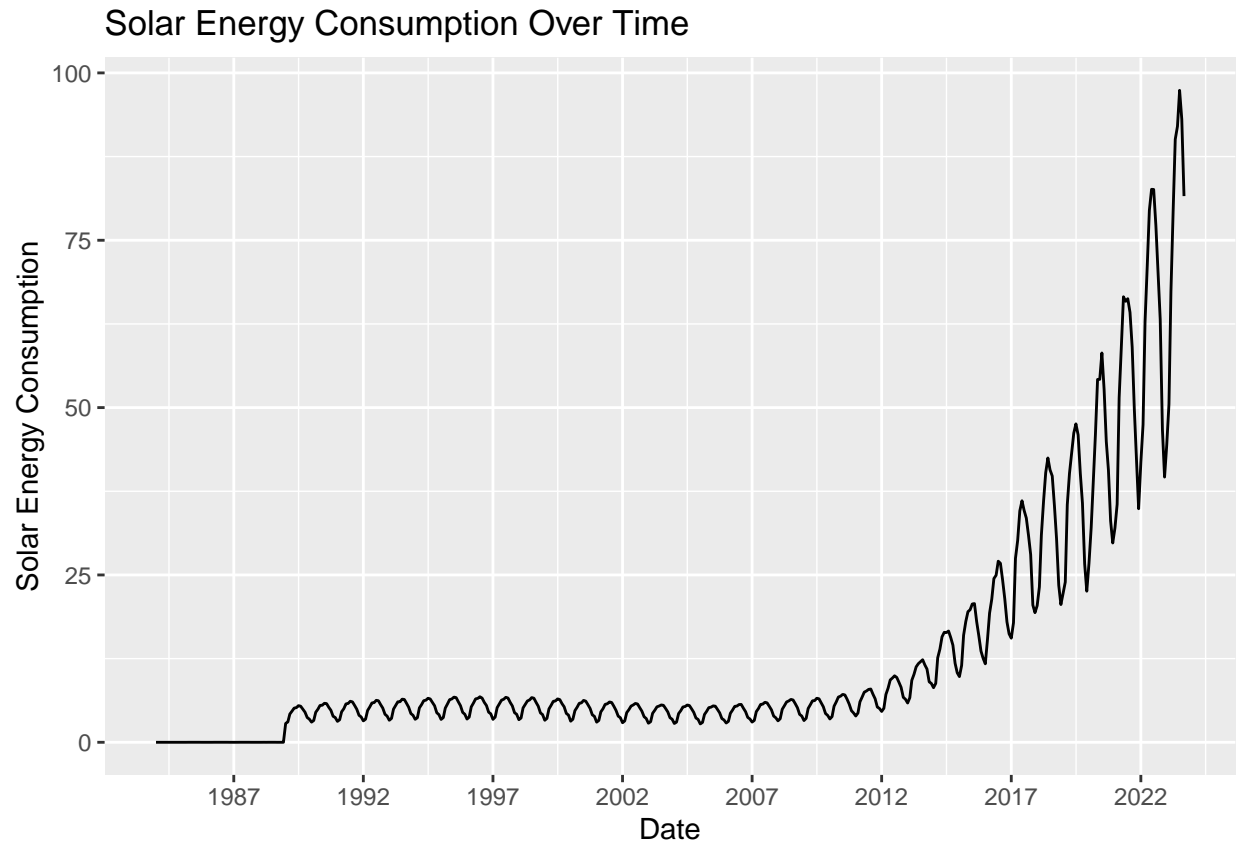
#Plot for Solar Energy Consumption
ggplot(energy_data, aes(x = Date, y = `Solar Energy Consumption`)) +
  geom_line() +
  ylab("Solar Energy Consumption") +
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +
  ggtitle("Solar Energy Consumption Over Time")

```

```

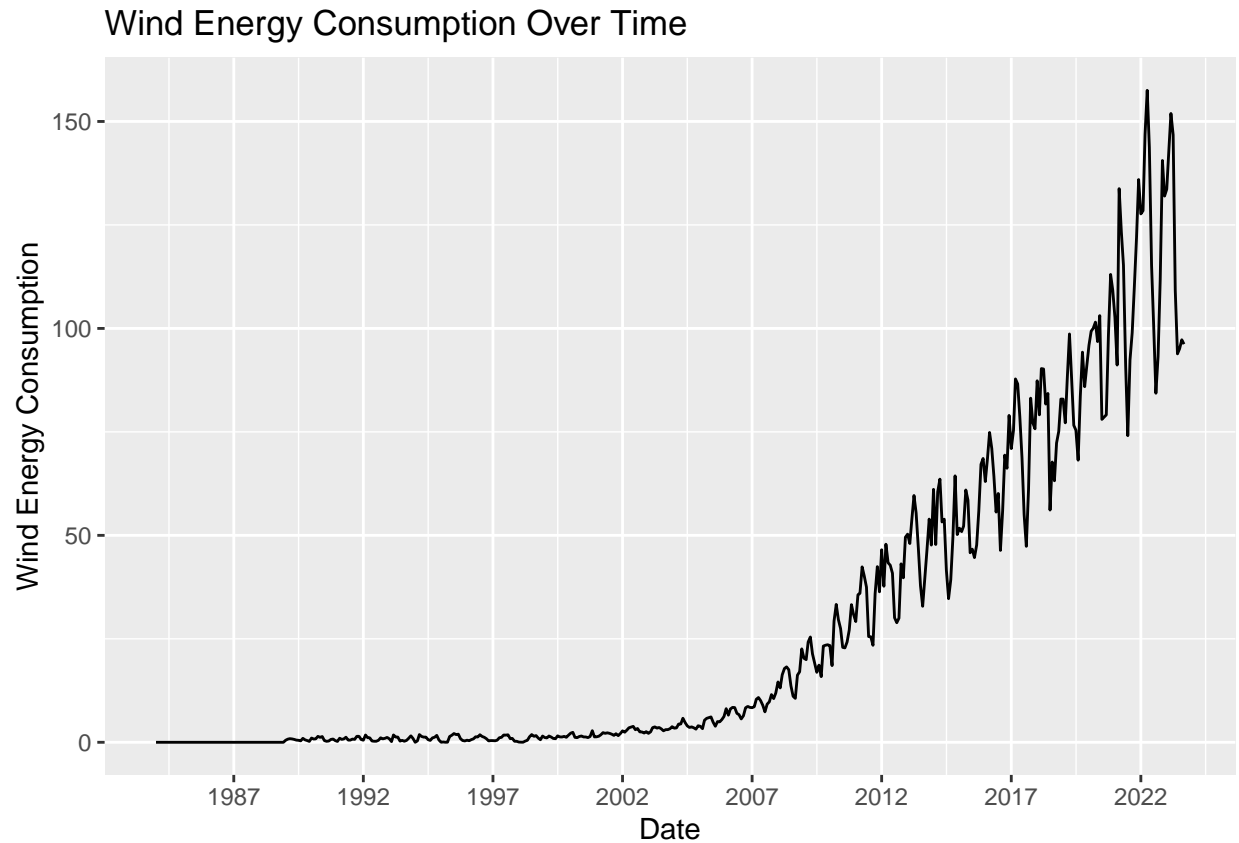
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.

```



```
#Plot for Wind Energy Consumption  
ggplot(energy_data, aes(x = Date, y = `Wind Energy Consumption`)) +  
  geom_line() +  
  ylab("Wind Energy Consumption") +  
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +  
  ggtitle("Wind Energy Consumption Over Time")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting  
## to continuous.
```

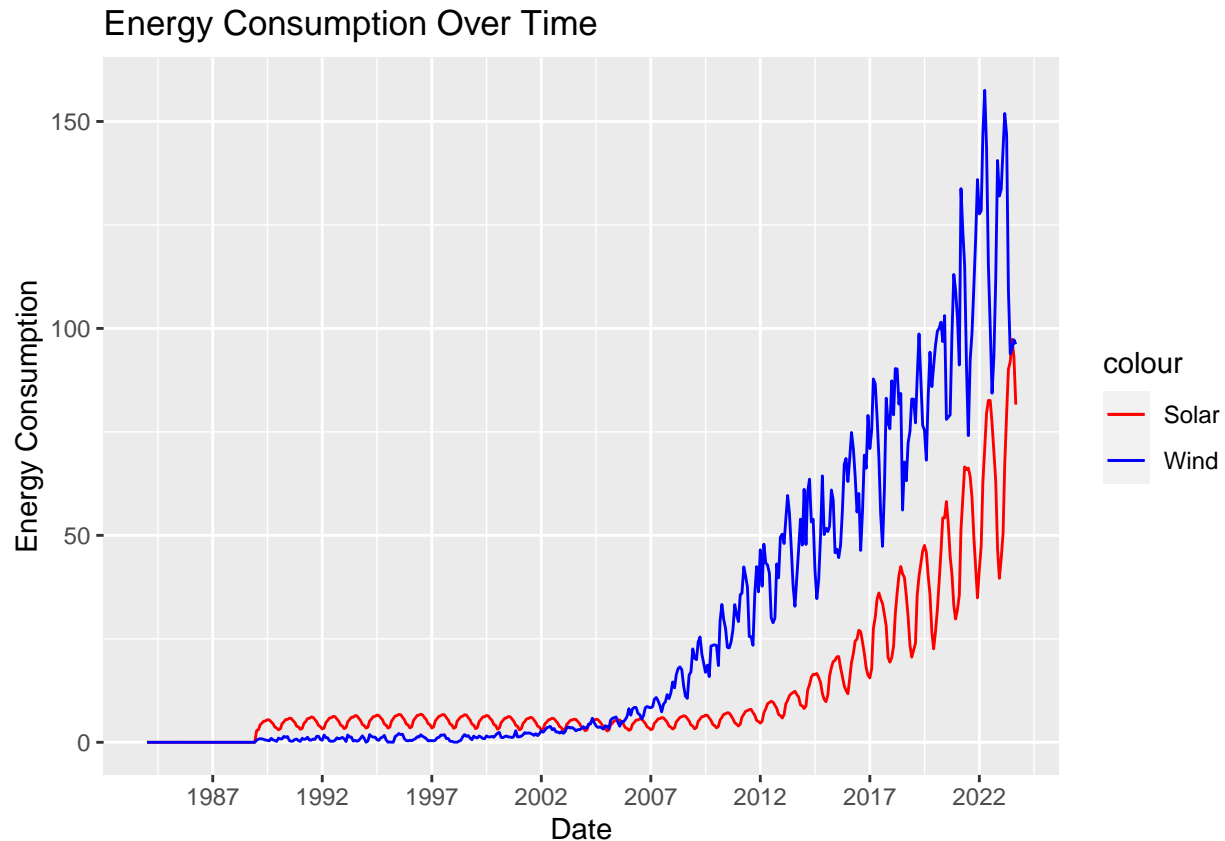


Q3

Now plot both series in the same graph, also using `ggplot()`. Use function `scale_color_manual()` to manually add a legend to `ggplot`. Make the solar energy consumption red and wind energy consumption blue. Add informative name to the y axis using `ylab("Energy Consumption")`. And use function `scale_x_date()` to set x axis breaks every 5 years.

```
ggplot(energy_data, aes(x=Date))+
  geom_line(aes(y=`Solar Energy Consumption`, color="Solar"))+
  geom_line(aes(y=`Wind Energy Consumption`, color="Wind"))+
  scale_color_manual(values=c("Solar"="red", "Wind"="blue"))+
  ylab("Energy Consumption")+
  scale_x_date(date_breaks="5 years", date_labels="%Y")+
  ggtitle("Energy Consumption Over Time")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```



Decomposing the time series

The stats package has a function called `decompose()`. This function only take time series object. As the name says the decompose function will decompose your time series into three components: trend, seasonal and random. This is similar to what we did in the previous script, but in a more automated way. The random component is the time series without seasonal and trend component.

Additional info on `decompose()`.

- 1) You have two options: alternative and multiplicative. Multiplicative models exhibit a change in frequency over time.
- 2) The trend is not a straight line because it uses a moving average method to detect trend.
- 3) The seasonal component of the time series is found by subtracting the trend component from the original data then grouping the results by month and averaging them.
- 4) The random component, also referred to as the noise component, is composed of all the leftover signal which is not explained by the combination of the trend and seasonal component.

Q4

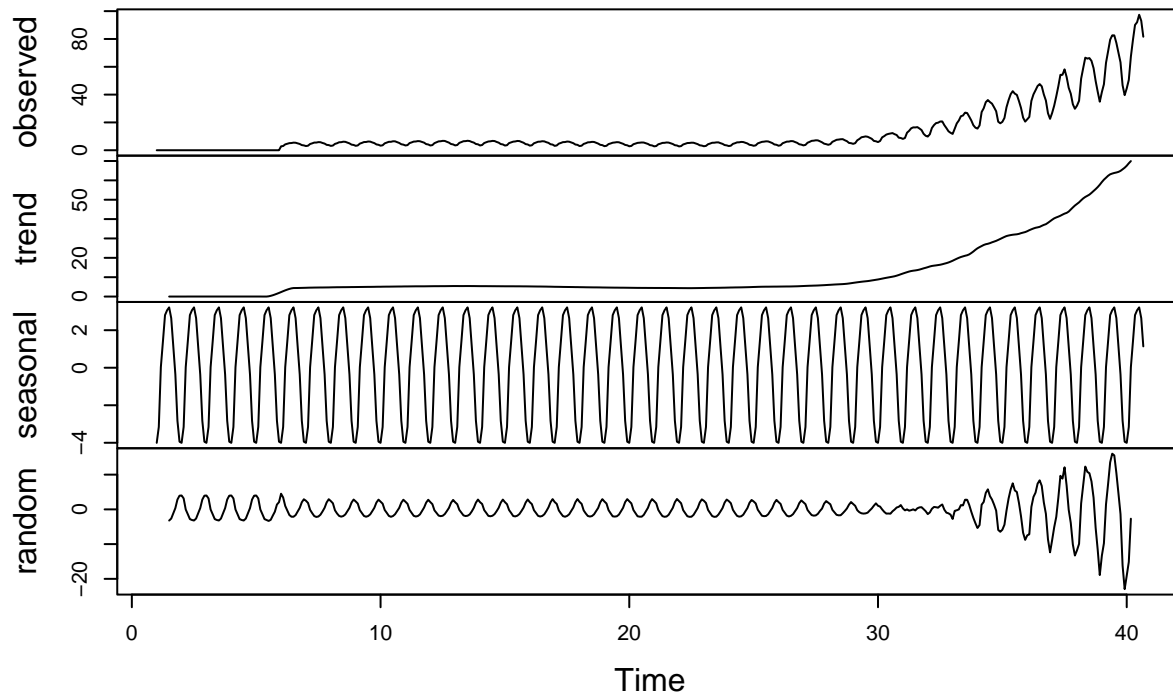
Transform wind and solar series into a time series object and apply the decompose function on them using the additive option, i.e., `decompose(ts_data, type = "additive")`. What can you say about the trend component? What about the random component? Does the random component look random? Or does it appear to still have some seasonality on it?

```
#Convert "Solar Energy Consumption" and "Wind Energy Consumption" to time series objects

solarTs<- ts(energy_data$`Solar Energy Consumption`,frequency=12)
SolarDecomposed <-decompose(solarTs,type ="additive")

plot(SolarDecomposed)
```

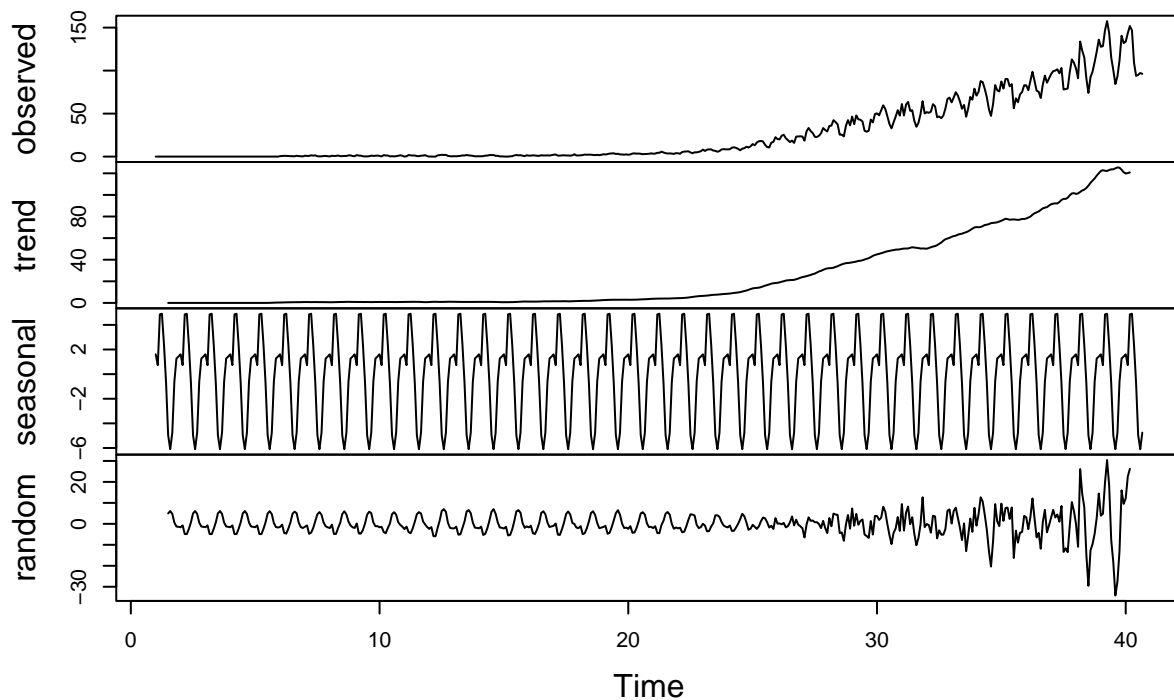
Decomposition of additive time series



```
windTs <-ts(energy_data$`Wind Energy Consumption`,frequency=12)
WindDecomposed <-decompose(windTs,type="additive")

plot(WindDecomposed)
```


Decomposition of additive time series



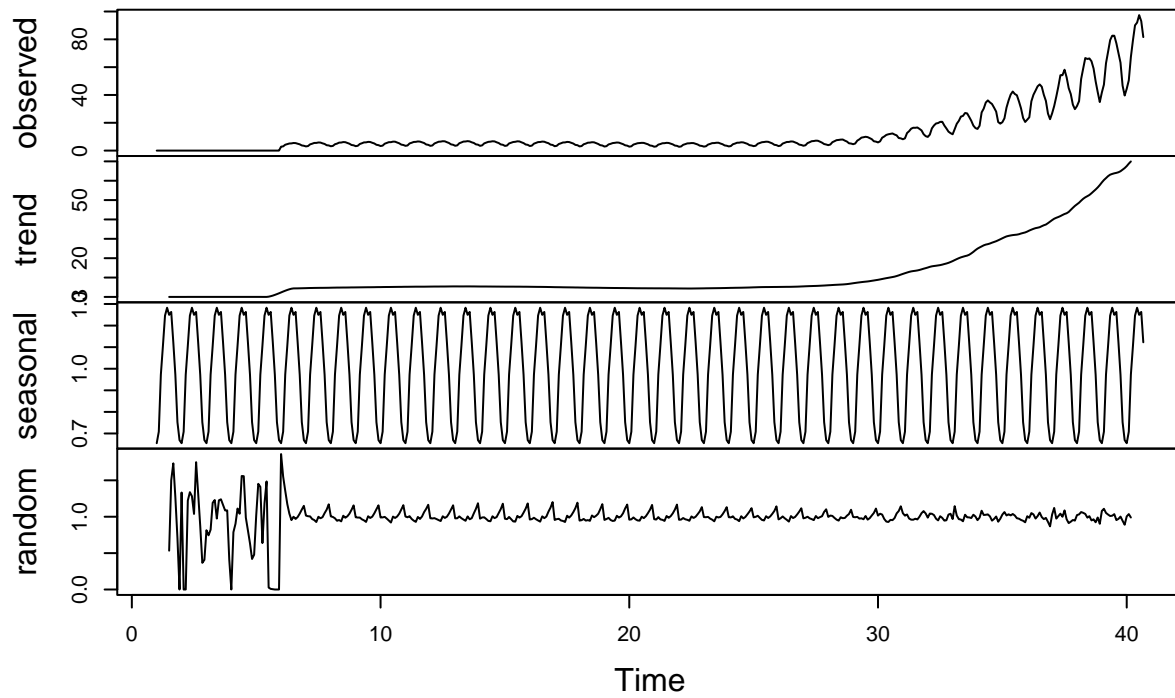
> Answer: For solar data and wind data they both demonstrate an exponential positive trend which peaks towards the end of the data set. For both data sets though the fluctuations only occur at a single time period in the data. While the rest remain symmetric suggesting that there is still some seasonality in the data set that wasn't yet removed.

Q5

Use the `decompose` function again but now change the type of the seasonal component from additive to multiplicative. What happened to the random component this time?

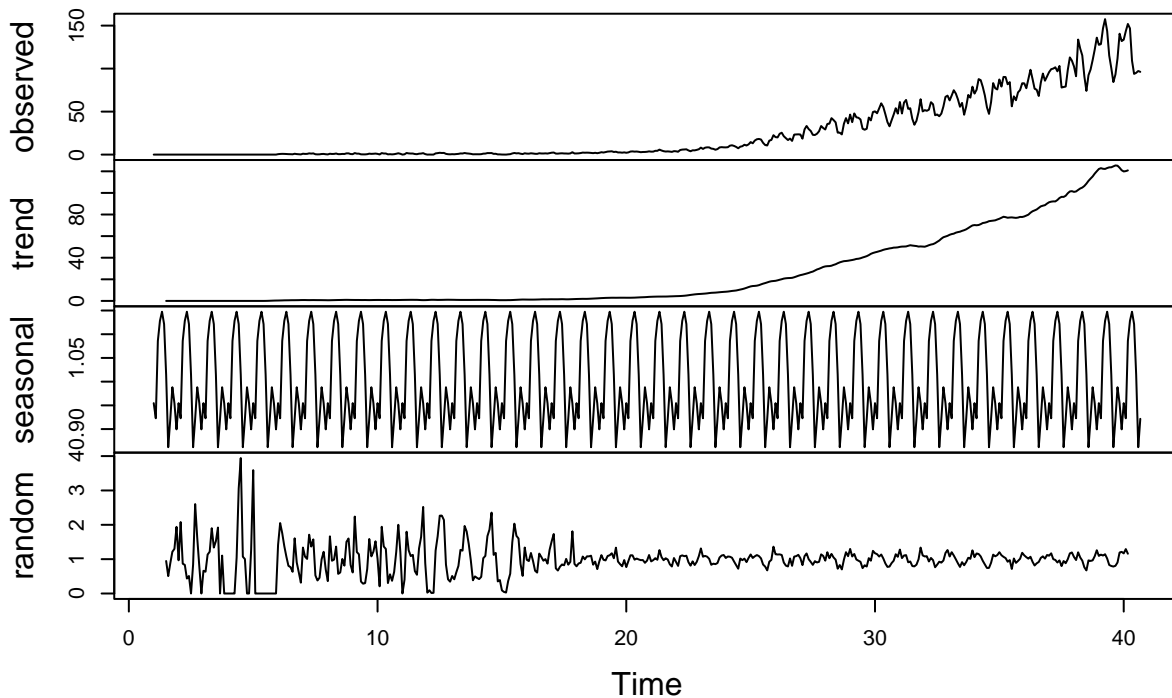
```
SolarDecomposed2 <- decompose(solarTs, type="multiplicative")
plot(SolarDecomposed2)
```

Decomposition of multiplicative time series



```
WindDecomposed2 <-decompose(windTs, type="multiplicative")  
plot(WindDecomposed2)
```

Decomposition of multiplicative time series



Answer: For the solar data we see some apparent randomness towards the beginning and end of the dataset. While the middle still displays seasonality. As for the wind data a significant portion shows some erratic fluctuations displaying that we might have captured the random component. But there is also a slight seasonality that should still be investigated further.

Q6

When fitting a model to this data, do you think you need all the historical data? Think about the data from the 90s and early 20s. Are there any information from those years we might need to forecast the next six months of Solar and/or Wind consumption. Explain your response.

Answer: Not necessarily, at times weather events that occurred in those years could be what is messing up the data. When looking at this data we see that it is not starting to change a fluctuate even more than before.

Q7

Create a new time series object where historical data starts on January 2012. Hint: use `filter()` function so that you don't need to point to row numbers, i.e. `filter(yyyy, year(Date) >= 2012)`. Apply the `decompose` function `type=additive` to this new time series. Comment the results. Does the random component look random? Think about our discussion in class about seasonal components that depends on the level of the series.

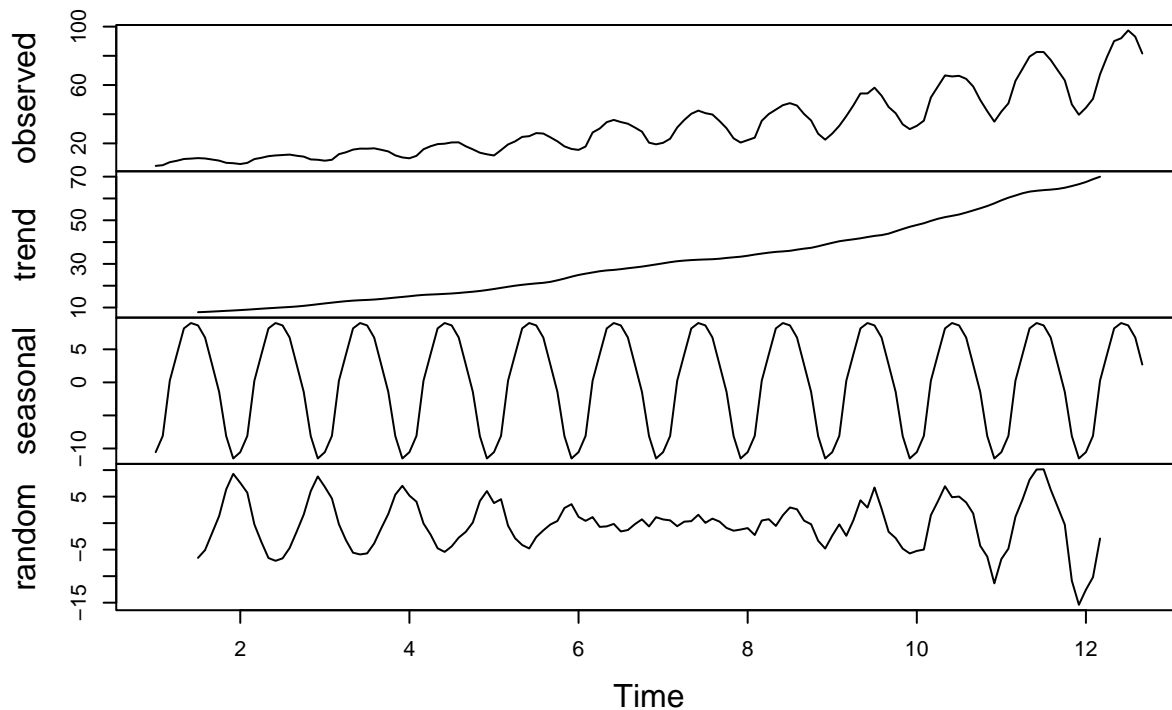
```

#Filter to start at January 2012
energydataNEW<-energy_data %>%
  filter(year(Date) >= 2012)

#Convert filtered data to time series object
#Apply decompose function with additive option
solarTsNEW<- ts(energydataNEW$`Solar Energy Consumption`,frequency=12)
solardecomposed2012 <- decompose(solarTsNEW, type="additive")
plot(solardecomposed2012)

```

Decomposition of additive time series

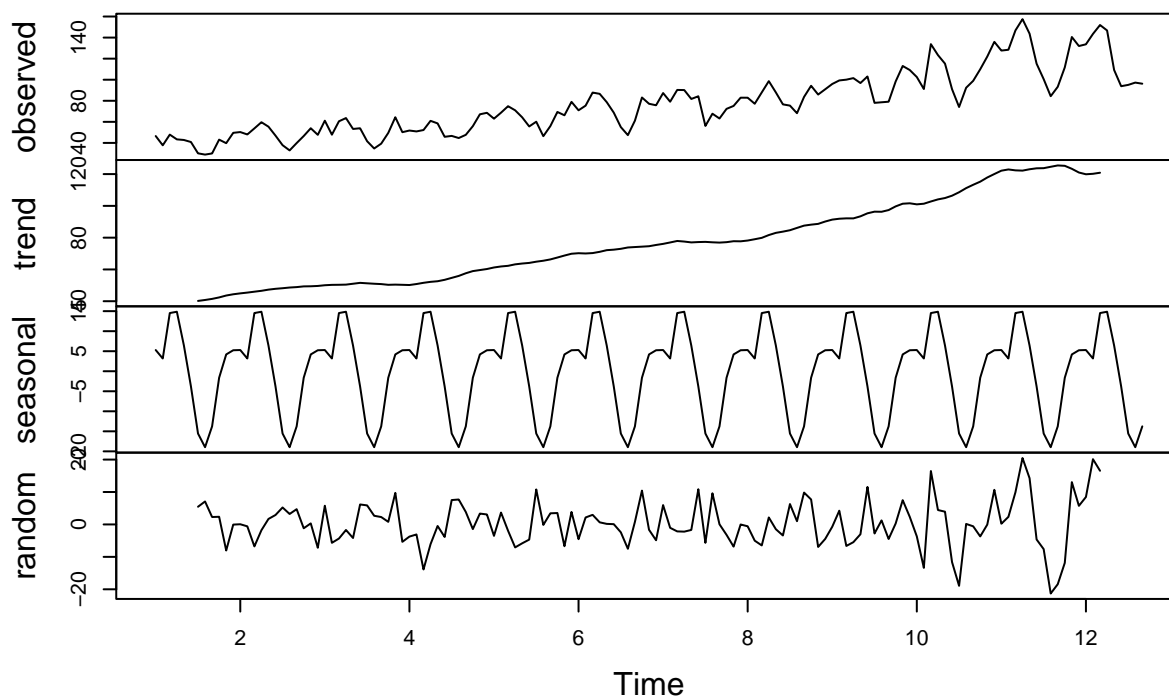


```

windTsNEW<- ts(energydataNEW$`Wind Energy Consumption`, frequency=12)
winddecomposed2012<- decompose(windTsNEW,type="additive")
plot(winddecomposed2012)

```

Decomposition of additive time series



Answer: Solar data demonstrates more randomness than previous plots, particularly the middle section, but still displays some seasonality. While, Wind data now displays the randomness that we were looking for. Though the levels are different this difference is what helped us remove the seasonality suggesting that what we needed to remove was within the years that we removed.

Identify and Remove outliers

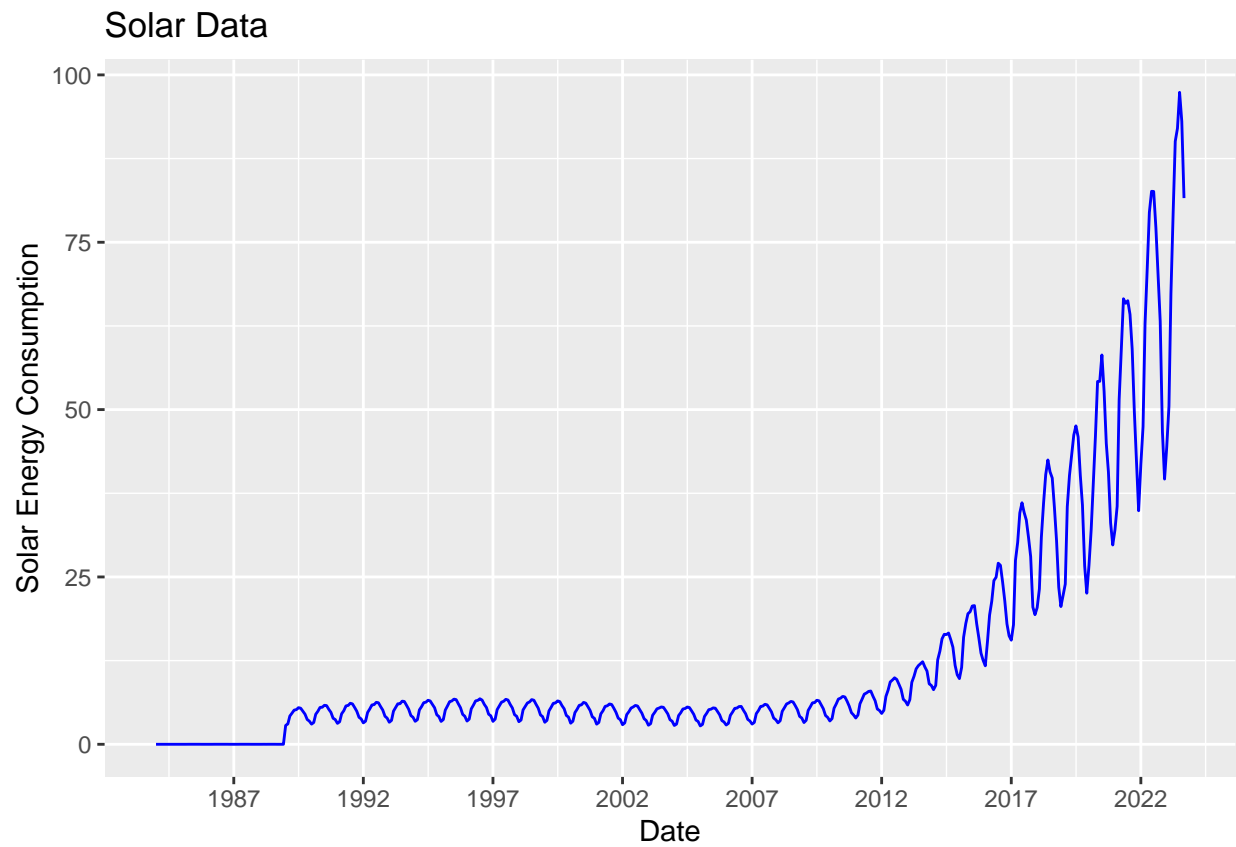
Q8

Apply the `tsclean()` to both series from Q7. Did the function remove any outliers from the series? Hint: Use `autoplot()` to check if there is difference between cleaned series and original series.

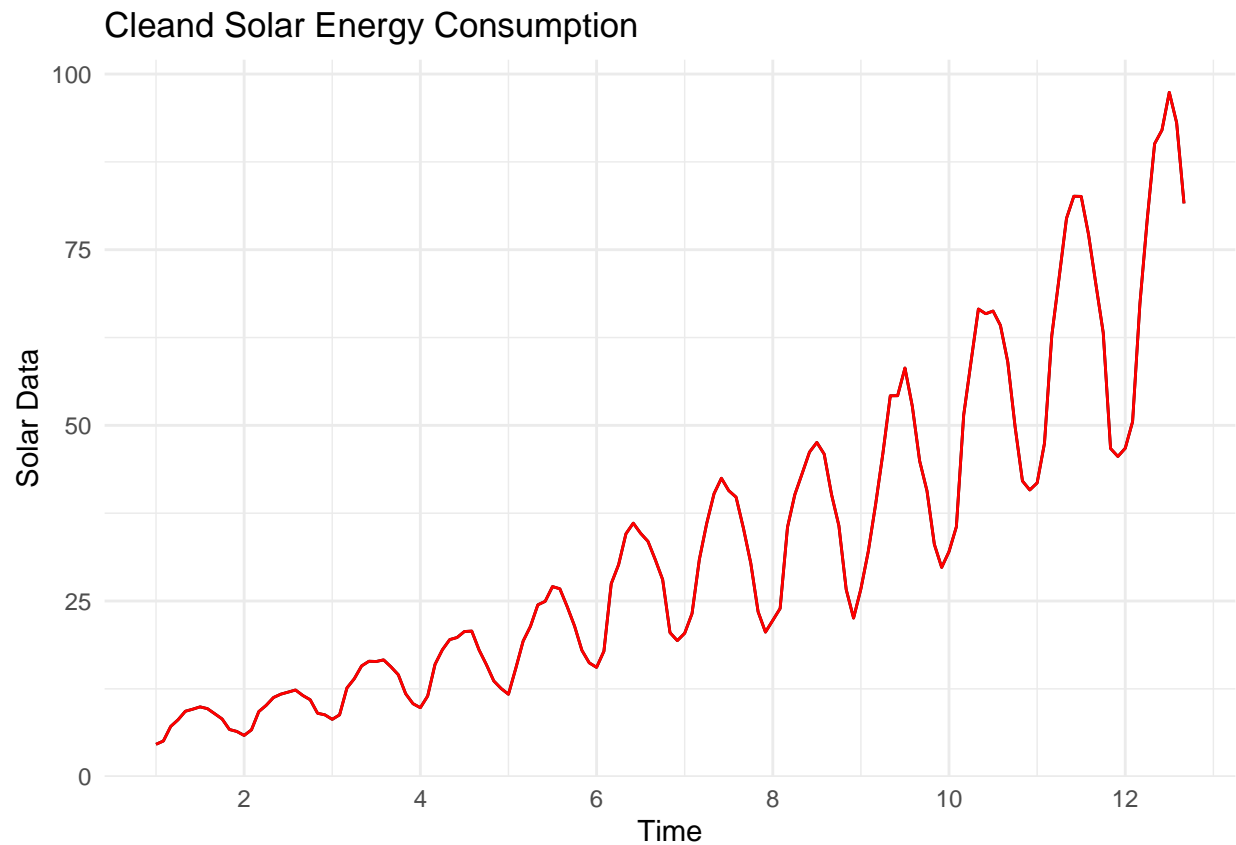
```
#Apply tsclean() to data
solar_cleaned <- tsclean(solarTsNEW)
wind_cleaned <- tsclean(windTsNEW)

ggplot(energy_data, aes(x=Date, y=`Solar Energy Consumption`))+
  geom_line(color="blue")+
  ylab("Solar Energy Consumption")+
  scale_x_date(date_breaks="5 years", date_labels="%Y")+
  ggtitle("Solar Data")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```

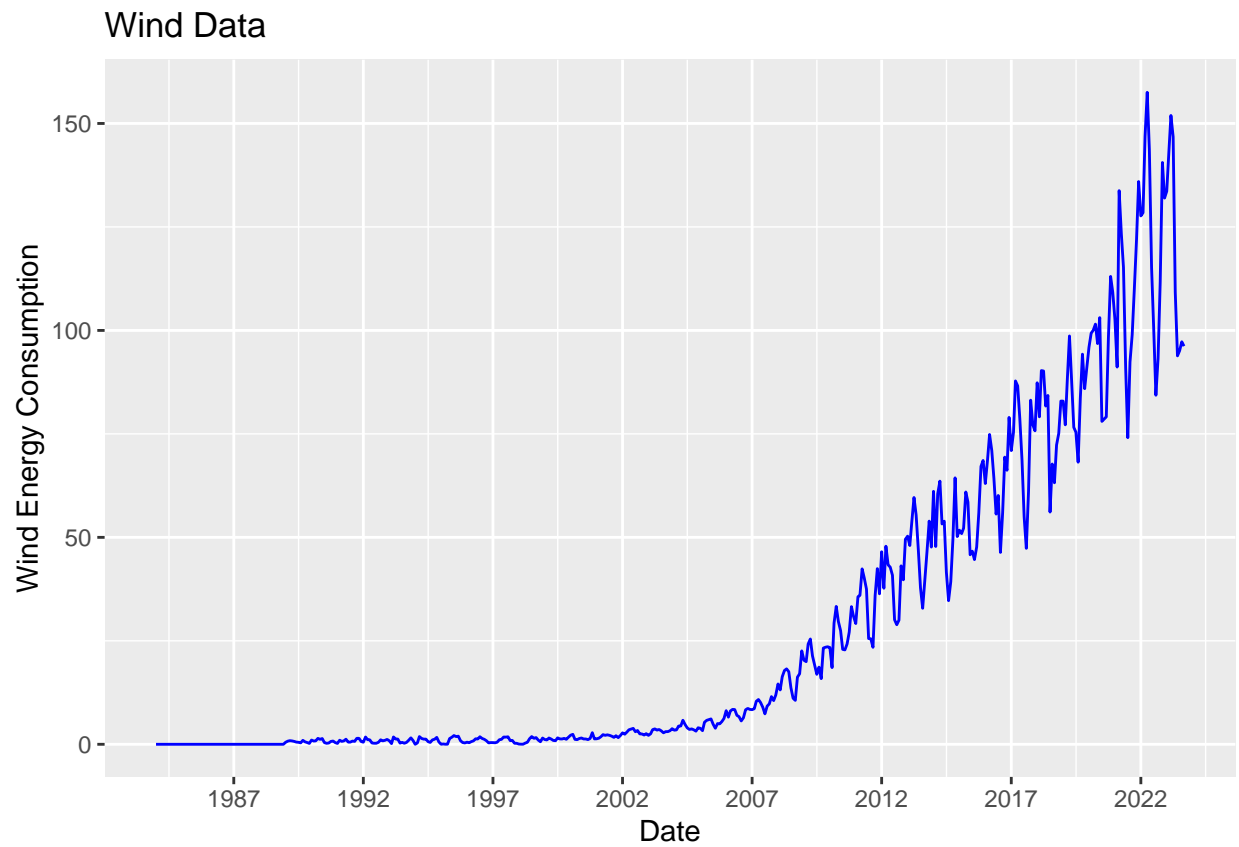


```
#plot using autoplot
autoplot(solar_cleaned)+
  geom_line(color="red")+
  ggtitle("Cleand Solar Energy Consumption") +
  theme_minimal() +
  ylab("Solar Data")
```

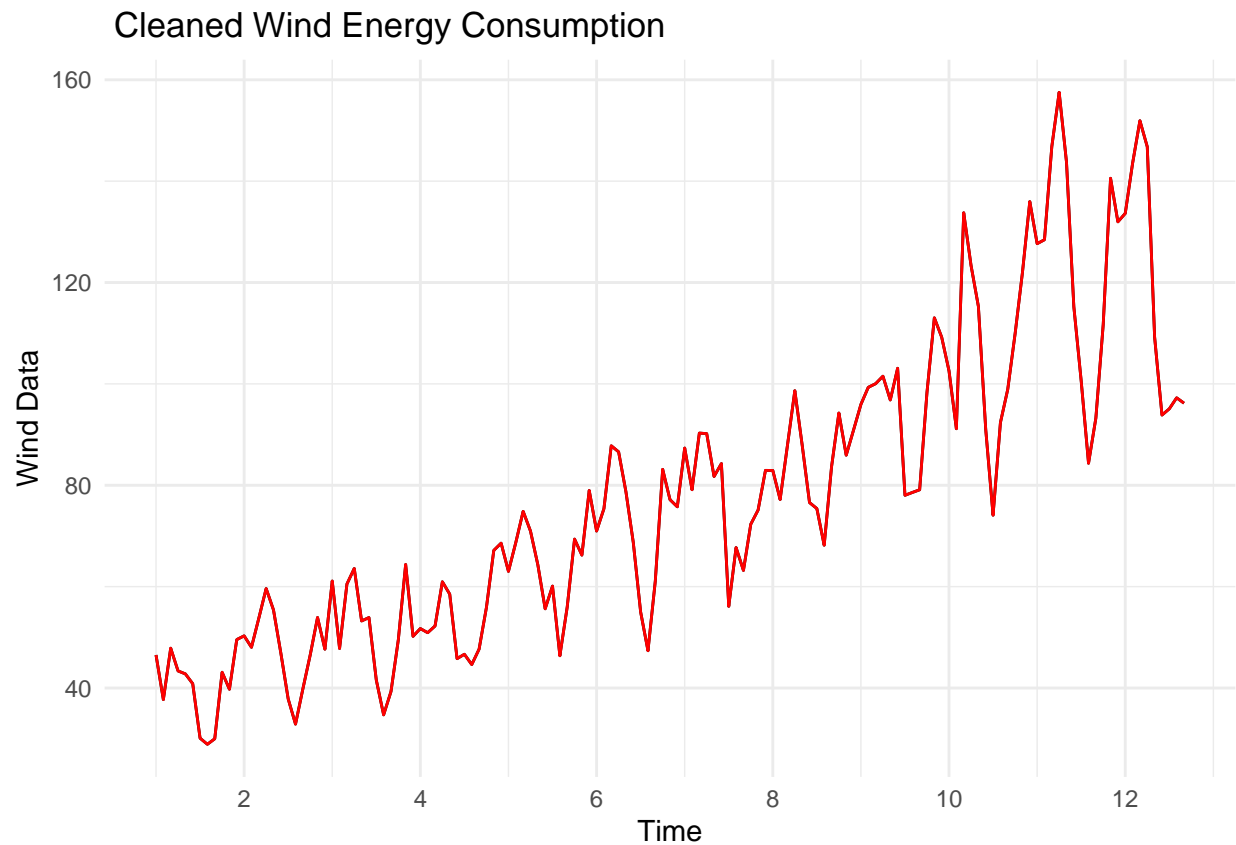


```
ggplot(energy_data, aes(x = Date, y = `Wind Energy Consumption`)) +  
  geom_line(color = "blue") +  
  ylab("Wind Energy Consumption") +  
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +  
  ggtitle("Wind Data")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting  
## to continuous.
```



```
#Plot using autoplot  
autoplot(wind_cleaned) +  
  geom_line(color = "red") +  
  ggtitle(" Cleaned Wind Energy Consumption") +  
  theme_minimal() +  
  ylab("Wind Data")
```

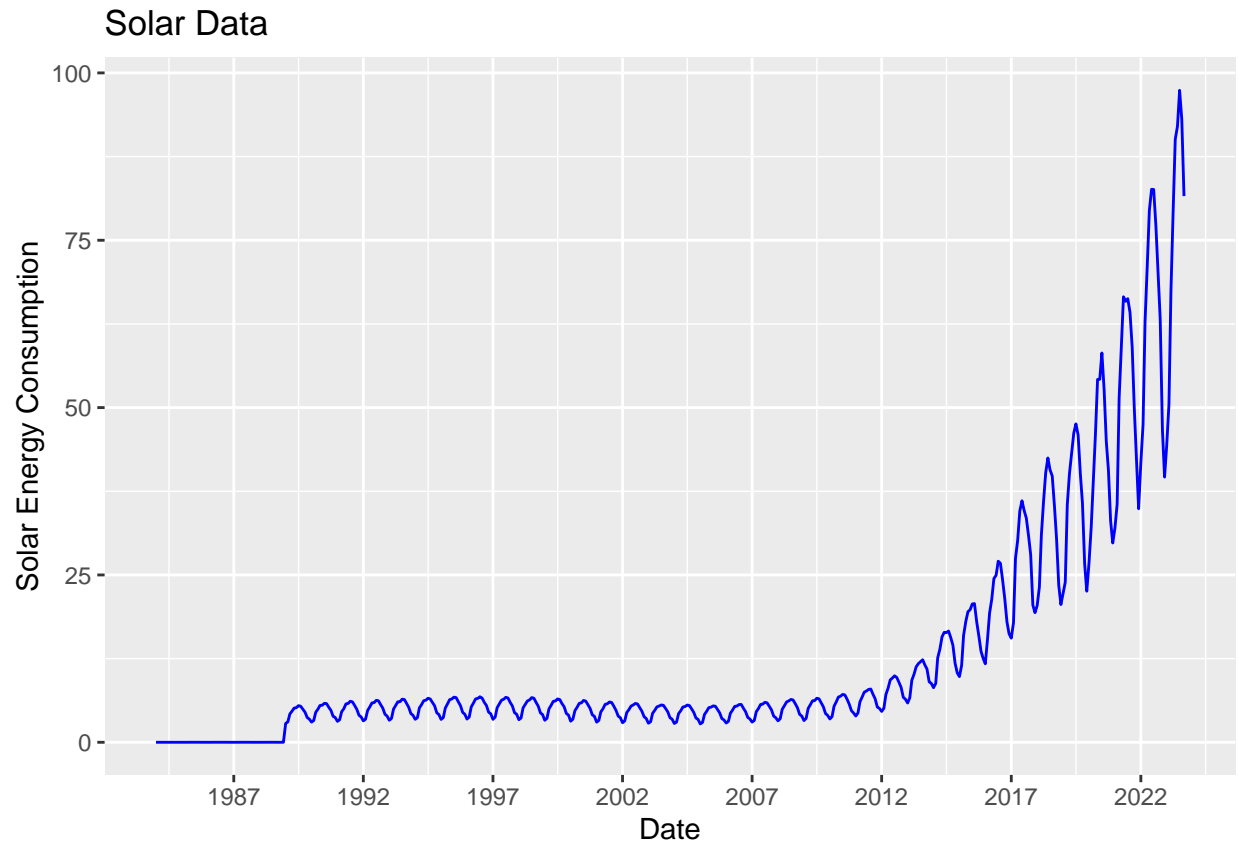



Q9

Redo number Q8 but now with the time series you created on Q7, i.e., the series starting in 2014. Using what `autoplot()` again what happened now? Did the function removed any outliers from the series?

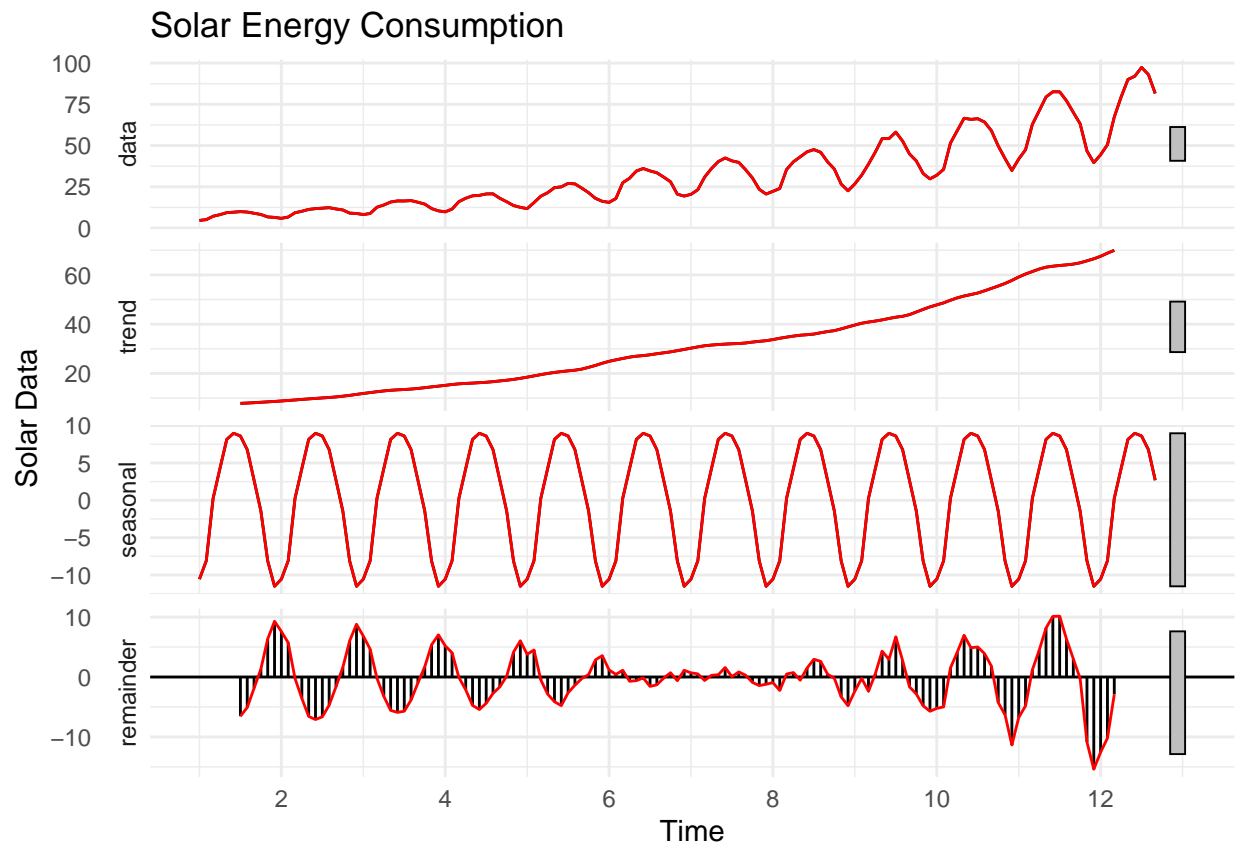
```
ggplot(energy_data, aes(x = Date, y = `Solar Energy Consumption`)) +
  geom_line(color = "blue") +
  ylab("Solar Energy Consumption") +
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +
  ggtitle("Solar Data")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```



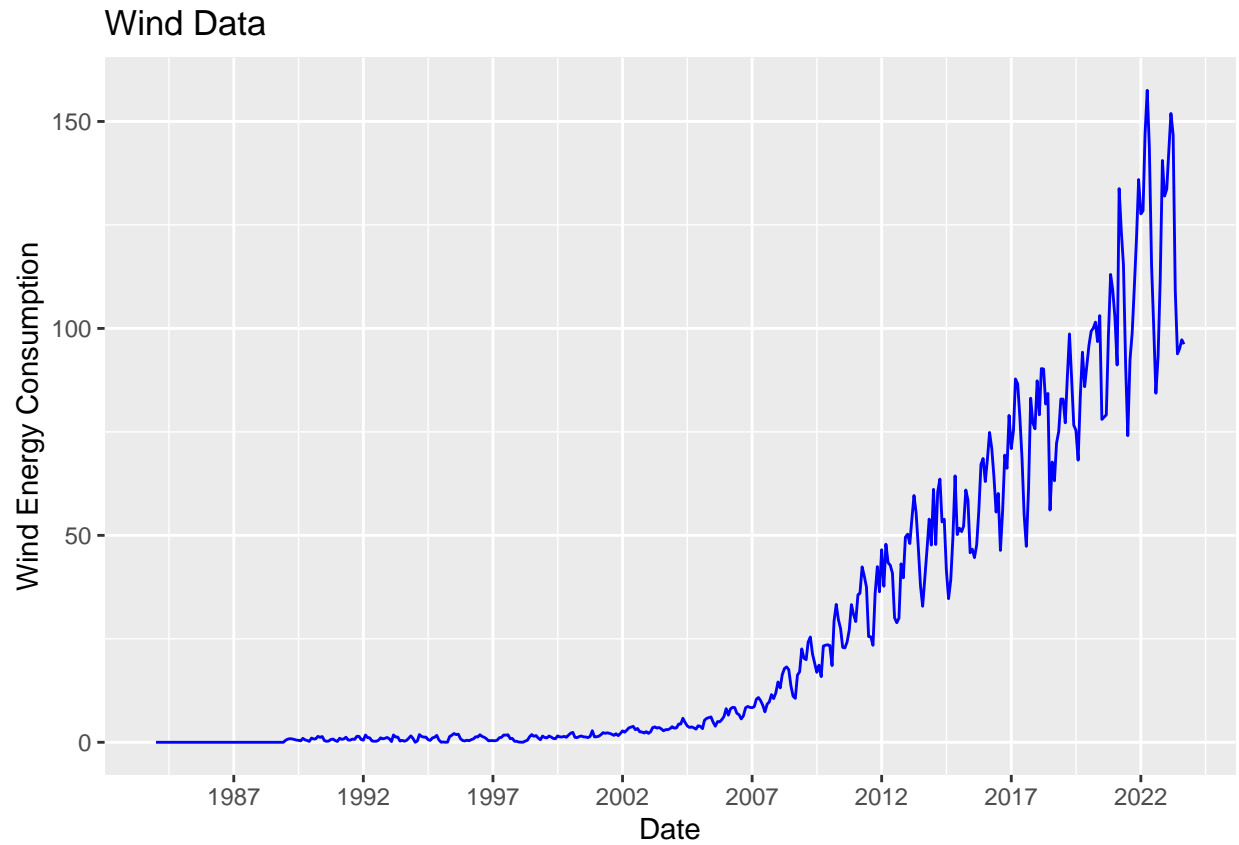
```
#Plot using autoplot  
autoplot(solardecomposed2012) +  
  geom_line(color = "red") +  
  ggtitle("Solar Energy Consumption") +  
  theme_minimal() +  
  ylab("Solar Data")
```

```
## Warning: Removed 6 rows containing missing values ('geom_line()').
```



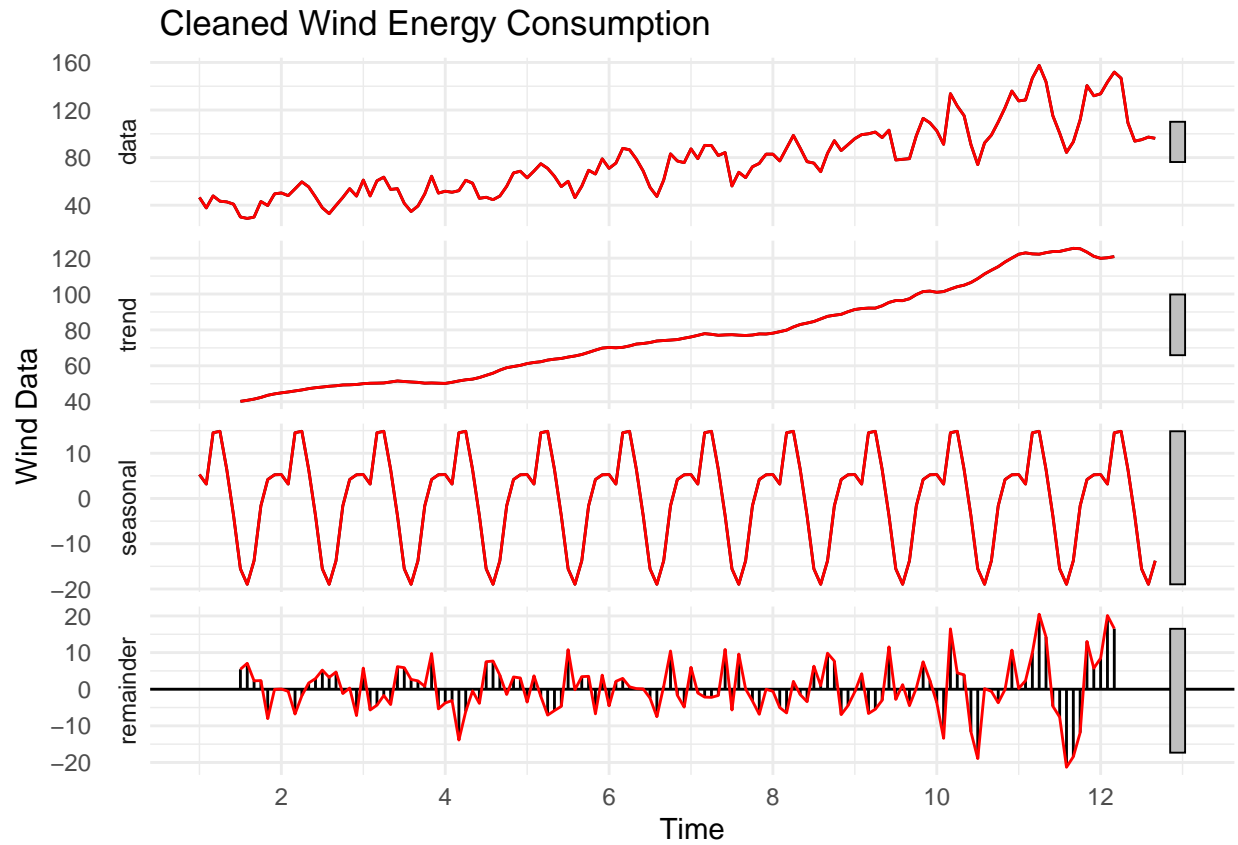
```
ggplot(energy_data, aes(x = Date, y = `Wind Energy Consumption`)) +
  geom_line(color = "blue") +
  ylab("Wind Energy Consumption") +
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +
  ggtitle("Wind Data")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```



```
#Plot using autoplot  
autoplot(windecomposed2012) +  
  geom_line(color = "red") +  
  ggtitle(" Cleaned Wind Energy Consumption") +  
  theme_minimal() +  
  ylab("Wind Data")
```

```
## Warning: Removed 6 rows containing missing values ('geom_line()').
```



Answer: We don't see much change between the two except that solar data have started to present a wave-like function with seasonality. While wind data became increasingly more random with the data that was filtered for less years.